

JANUARY 20, 2018



HEALTHCARE SERVICES

PRESENTED BY: YASH MITTAL
PRESENTED TO: EDWISOR.COM

CONTENTS

1. Problem Statement

2. Data used

3. Exploration of Numerical Variables

4. Exploration of Categorical Variables

5. Hypothesis Testing

5.1. CHI-SQUARE TEST

6. Data Pre-processing

7. Building Predictive Models

7.1. SGD Classifier

7.2. Multinomial Naïve Bayes

8. Conclusion

9. Recommendations

1. Problem Statement:

XYZ Health Services is a top ranked health care provider in USA with stellar credentials and provides high quality-care with focus on end-to-end health care services. The health care services range from basic medical diagnostics to critical emergency services. The provider follows a ticketing system for all the telephonic calls received across all the departments. Calls to the provider can be for new appointment, cancellation, lab queries, medical refills, insurance related, general doctor advice etc. The tickets have the details of summary of the call and description of the calls written by various staff members with no standard text guidelines.

The challenge is, based on the text in the Summary and description of the call, the ticket is to be classified to appropriate category (out of 5 categories) and subcategories (out of 20 sub categories).

It is a text classification problem.

2. Data Used:

We are provided with a csv file which contain 57280 rows and 7 columns. These 7 columns are:

- fileid
- summary
- data
- categories
- sub_categories
- previous_appointment
- id

categories and sub_categories are our target variables which we have to predict via other variables. fileid and id are used for numbering of the data in csv, and these are of no use in predicting our target variables. We will remove these 2 variables and then we will be working with 5 variables. We are now only left with 3 variables that are summary, data and previous_appointment which can be used for prediction.

3. Exploration of Numerical variables:

We have no numerical variables on which we need to do analysis as the numerical variables fileid and id are already left behind.

4. Exploration of Categorical variables:

We have 3 categorical variables in data:

- previous_appointment
- category
- sub_category

In previous_appointment we get counts of 'no' and 'yes' as:

```
no      57083
yes      195
Name: previous_appointment, dtype: int64
```

There are 2 blank entries in previous_appointment variable.


After removal of all incomplete rows we get value counts for category and sub category as:

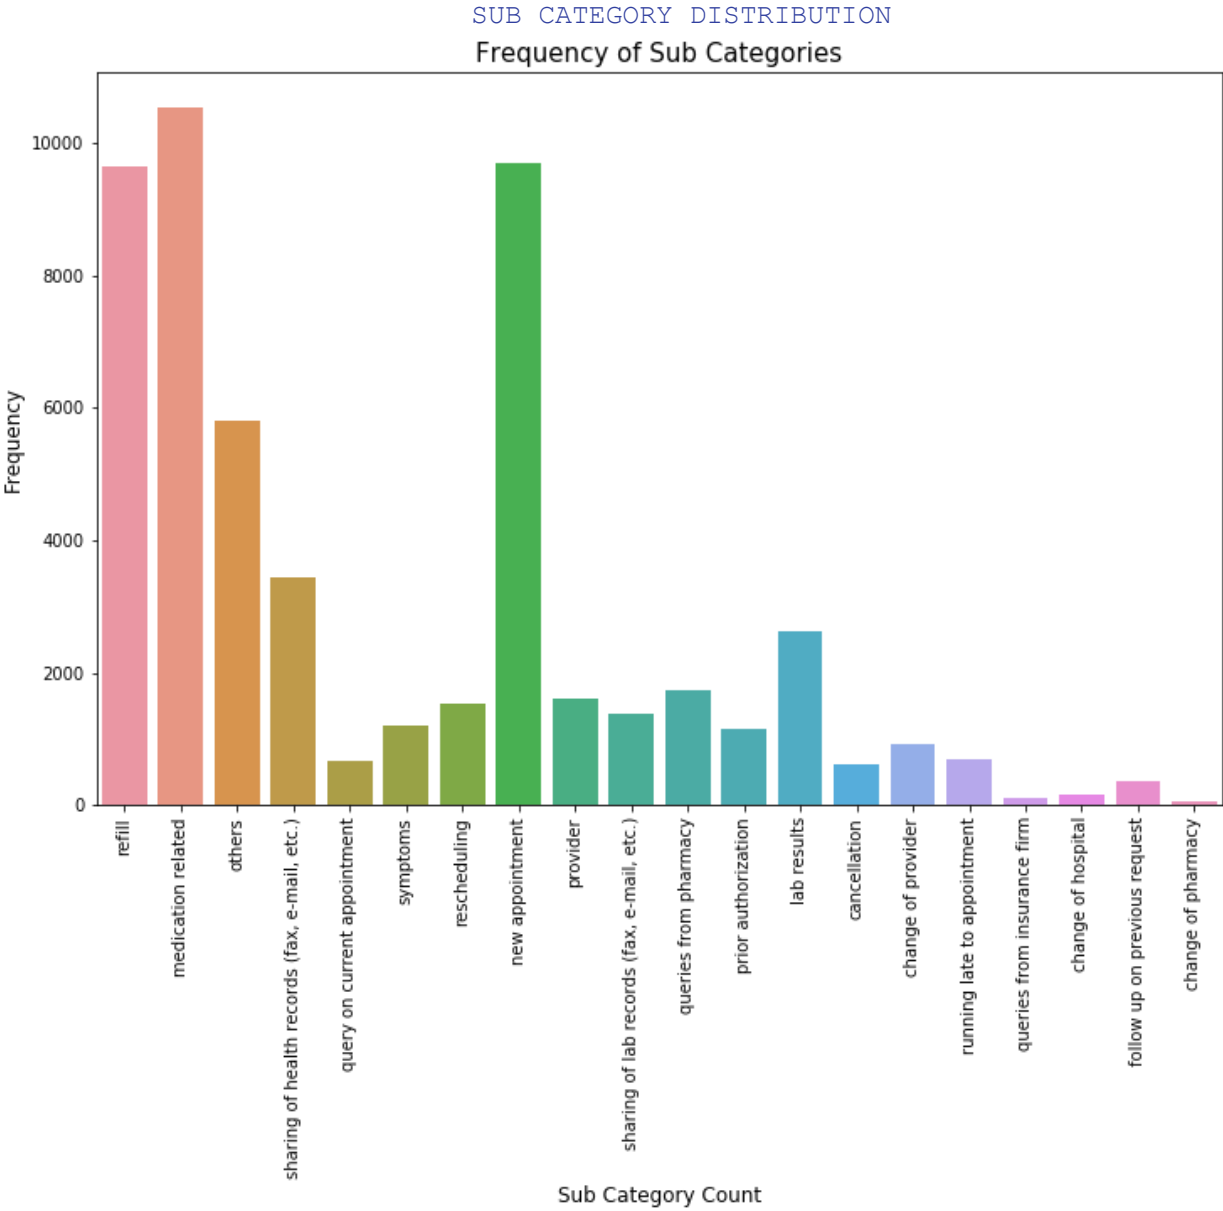
```
prescription      14500
appointments      12960
ask_a_doctor      11744
miscellaneous      10462
lab                4246
Name: categories, dtype: int64
```

```
medication related      10547
new appointment         9713
refill                  9665
others                  5796
sharing of health records (fax, e-mail, etc.)  3435
lab results             2628
queries from pharmacy   1722
provider                1608
rescheduling            1520
sharing of lab records (fax, e-mail, etc.)    1386
symptoms                1197
prior authorization     1155
change of provider       928
running late to appointment  694
```


query on current appointment	652
cancellation	613
follow up on previous request	350
change of hospital	142
queries from insurance firm	107
change of pharmacy	54
Name: sub_categories, dtype: int64	

CATEGORY DISTRIBUTION

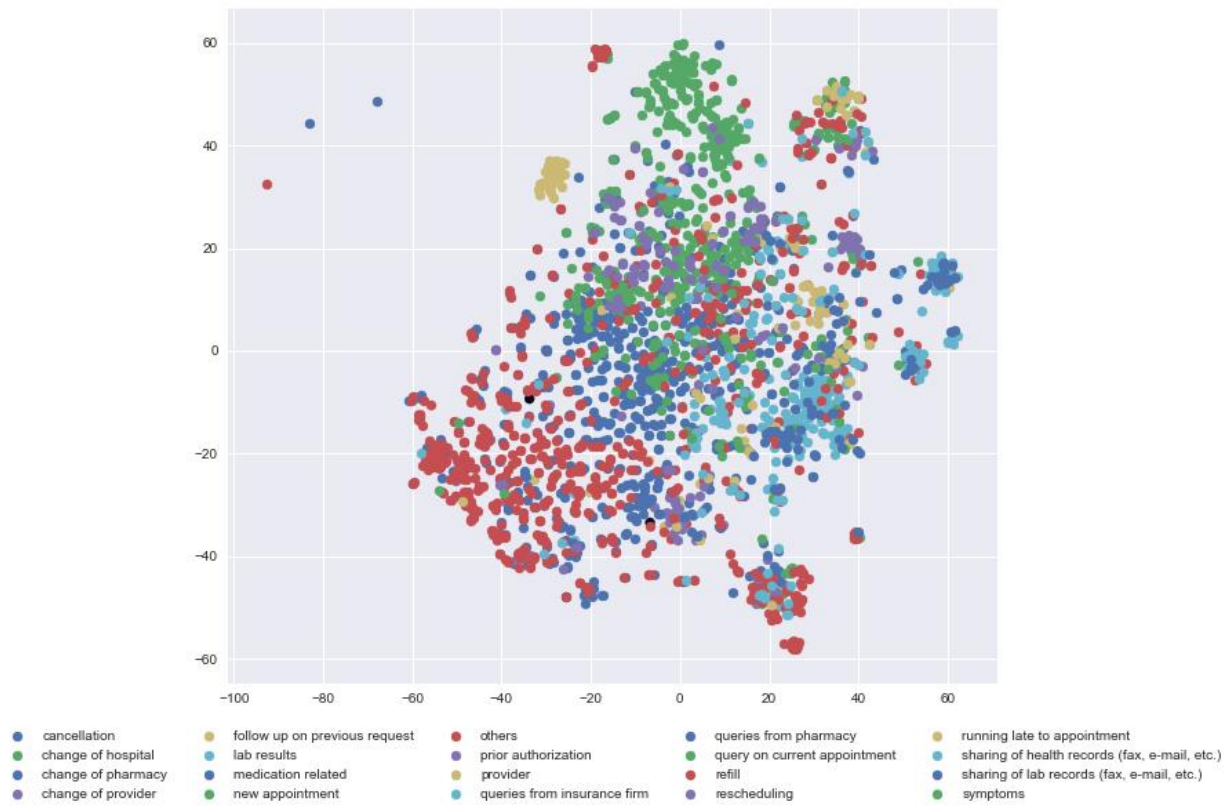




We can also visualize this high dimensional data of categories and sub categories using t-SNE.



Sub Categories t-SNE Visualization



5 . Hypothesis Testing:

The P-value approach involves determining "likely" or "unlikely" by determining the probability — assuming the null hypothesis were true — of observing a more extreme test statistic in the direction of the alternative hypothesis than the one observed. If the P-value is small, say less than (or equal to) α , then it is "unlikely." And, if the P-value is large, say more than α , then it is "likely."

If the P-value is less than (or equal to) α , then the null hypothesis is rejected in favour of the alternative hypothesis. And, if the P-value is greater than α , then the null hypothesis is not rejected.

5.1 Chi-Square Test:

We will use chi-square test to find the dependency of categorical variables on each other.

- **Hypothesis testing on previous_appointment and categories:**
We will assume that they are independent and when we perform chi square test by assuming alpha as 0.05 we get value of p as 7.1869681817705634e-16 which is quite smaller to alpha. So our hypothesis failed and they are dependent on each other
- **Hypothesis testing on categories and sub_categories:**
We will assume that they are independent and when we perform chi square test by assuming alpha as 0.05 we get value of p as 0 which is too smaller to alpha. So our hypothesis failed and they are strongly dependent on each other.

6. Data Pre-processing:

DATA variable contains a lot of junk in it which needs to be cleared before putting the data into various machine learning models. Example of a cell is:

```
'{\rtf1\ansi\ftnbj{\fonttbl{\f0 \fswiss Arial;}}{\colortbl ;\red255\green255\blue255 ;\red0\green0\blue255 ;\red0\green0\blue0 ;\red0\green0\blue255 ;\red0\green128\blue0 ;}{\stylesheet{\f0\fs20\cf3\cb1 Normal;}{\cs1\additive\cf3\cb1 Default Paragraph Font;}}\margl1440\margr1440\margt540\margb1440\headery540\footery720\formshade\sectd\marglsxn1440\margrsxn1440\margtsxn540\margbsxn1440\headery540\footery720\sbkpage\pgncont\plain\plain\fs20\pard\plain\fs20\cf0\fs24\sscharaux1\b Phone Note \fs20\b0\par\fs24\b\par CALL FROM PHARMACY\par \fs20 Caller Name: \b0xxxx-xxxxar\b Caller: \b0 right source\par\b\par Reason for Call: \b0 Details: 866-4862668\par Requesting verbal bc original rx was an ink blurr.\par\par Plegridy\par\b Call Taken by: \b0 xxxx-xxxx , January 5, 2015 2:12 PM\par\b Follow-up Details: \b0 2nd call concerning verbal order for Plegridy\par\b Follow-up by: \b0 Paulette Lee, January 5, 2015 4:25 PM\par\b Additional Follow-up Details: \b0 3rd call concerning verbal order for Plegridy, pt was to start taking it today. \par\b Additional Follow-up by: \b0 Paulette Lee, January 6, 2015 8:35 AM\par\fs24\b\par Clinical List Changes\par\fs20\par Medications Updated:\par\b0 Added new medication of PLEGRIDY STARTER PACK 63 & 94 MCG/0.5ML SOPN (PEGINTERFERON BETA-1A) Inject 63mcg SQ Day 1, then Inject 94 mcg Day 15, then start maintenance dose 125 mcg thereafter q 14 days - Signed\par Removed medication of AVONEX PEN 30 MCG/0.5ML KIT (INTERFERON BETA-1A) inject once a week via Avonex Pen\par Added new medication of PLEGRIDY 125 MCG/0.5ML SOPN (PEGINTERFERON BETA-1A) Inject 125 mcg SQ q 14 days - Signed\par Rx of PLEGRIDY STARTER PACK 63 & 94 MCG/0.5ML SOPN (PEGINTERFERON BETA-1A) Inject 63mcg SQ Day 1, then Inject 94 mcg Day 15, then start maintenance dose 125 mcg thereafter q 14 days; #1 Package x 0; Signed; Entered by: Lisa Jones RN; Authorizedxxxx-xxxxWEST '

```

This contain various '\ ' signs with a lot of garbage data with them.

Following code is used to clean data from all the rows:

```
%%time
for i in range(0,df.shape[0]):
    print i
    df.iloc[i,2]=re.sub(r'\\..*? |\\..*;',',',df.iloc[i,2])
    df.iloc[i,2]=re.sub(r'xxxx|xxx|-',',',df.iloc[i,2])
    df.iloc[i,2]=re.sub(r'.+;}',',',df.iloc[i,2])
    df.iloc[i,2] = re.sub(r'[a-zA-Z]+\s+\d{1,2},\s+\d{4}\s+\d{1,2}:\d{2}\s+(AM|PM){0,1}','',df.iloc[i,2])

```

Data of dates is also removed since we don't need them while building our bag of words model. Now the clean data is saved in processed.csv file so that next time we use data we don't need to process that again.

There are also missing values in our data which needs to be removed.

```
df1.isnull().sum()
```

```
SUMMARY          3347
DATA              0
categories        0
sub_categories    0
previous_appointment  2
dtype: int64
```

There are 3347 null values in SUMMARY variable and 2 null values in previous_appointment.

```
df1=df1.dropna()
```

All null values will be dropped from the dataframe.

We will also convert all data of categories and sub_categories into lower case so that they doesn't get treated as different.

In some rows values of categories and sub categories are missing and that is filled with word 'junk'. Those rows will create problem while prediction so we will also drop those rows. After removing all unwanted rows we are now left with 53912 rows.

Now, we will make use of 2 columns DATA and SUMMARY and will make a single column total_data by merging them.

This total data will be cleaned now and values of previous_appointment variable will also be added into it. Data like punctuation marks are removed and only plain text is kept. Data is converted into lower case. All the extra whitespaces are removed.

We will also perform stemming there so that words like 'patient' and 'patients' are treated as same. The following code is doing that and making corpus for our upcoming model:


```

%%time
corpus=[]
fdf['clean_data']="
for i in range(0, fdf.shape[0]):
    print(i)
    review = re.sub('[^a-zA-Z]', ' ', fdf['total_data'][i])
    review = review.lower()
    review = review.split()
    ps = PorterStemmer()
    review = [ps.stem(word) for word in review if not word in set(stopwords.words('english'))]
    review = ' '.join(review)
    review=review+' '+fdf['previous_appointment'][i].lower()
    fdf['clean_data'][i]=review

```

This data is saved to file final.csv so that we don't need to do same processing again when coming back on work.

Corpus visualization using word cloud:



We can see some most commonly words in word cloud like call, patient, caller phone, follow, detail, etc.

7. Building Predictive Models:

7.1. SGD Classifier:

This estimator implements regularized linear models with stochastic gradient descent (SGD) learning: the gradient of the loss is estimated each sample at a time and the model is updated along the way with a decreasing strength schedule (aka learning rate). For best results using the default learning rate schedule, the data should have zero mean and unit variance.

This implementation works with data represented as dense or sparse arrays of floating point values for the features. The model it fits can be controlled with the loss parameter; by default, it fits a linear support vector machine (SVM).

I am creating pipeline for this using Countvectorizer and Tf-idf:

```
category_clf = Pipeline([('vect', CountVectorizer(ngram_range=(1,2))), ('tfidf', TfidfTransformer()),
('clf', SGDClassifier(loss = 'modified_huber', alpha=0.0001, penalty = 'l2', n_iter = 100, random_state = 0))])
```

The data is then divided into training and testing data in proportion of 80:20 using stratified sampling.

Classifier is trained and saved on harddisk:

```
with open('SGD_category.pkl', 'wb') as f:
    pickle.dump(category_clf, f)
```

We obtain accuracy of 82.5% through this classifier. In cross validation check we found that accuracy is around 81.7% which is quiet close to it.

Following is the metrices classification report:

	precision	recall	f1-score	support
appointments	0.88	0.86	0.87	2693
ask_a_doctor	0.76	0.78	0.77	2305
lab	0.85	0.76	0.80	820
miscellaneous	0.76	0.79	0.77	2099
prescription	0.85	0.87	0.86	2866
avg / total	0.82	0.82	0.82	10783

Now for predicting sub categories, categories will also play an important role so we will add categories column also to our corpus for predicting sub categories

and save merged column `clean_data_categories` with dataframe to `final_sub_categories.csv`.

For predicting sub categories again we split data into training and test data in proportion of 80:20. Classifier of sub category prediction is made:

```
sub_category_clf = Pipeline([('vect', CountVectorizer(ngram_range=(1,2))), ('tfidf', TfidfTransformer()),
                              ('clf', SGDClassifier(loss = 'modified_huber', alpha=1e-05, penalty = 'l2', n_iter = 100, random_state = 0))])
```

We train the model and save it with:

```
with open('SGD_sub_category.pkl', 'wb') as f:
    pickle.dump(sub_category_clf, f)
```

For predicting sub category I predict category first and then added that category to the corpus to predict sub category. Through this I achieved accuracy of 80.8% almost 81% on sub category prediction.

Metrics classification report for sub category prediction is:

	precision	recall	f1-score	support
cancellation	0.87	0.58	0.69	123
change of hospital	0.00	0.00	0.00	28
change of pharmacy	0.00	0.00	0.00	11
change of provider	0.59	0.59	0.59	186
follow up on previous request	0.55	0.17	0.26	70
lab results	0.84	0.90	0.87	526
medication related	0.80	0.86	0.83	2110
new appointment	0.83	0.90	0.86	1943
others	0.75	0.79	0.77	1159
prior authorization	0.78	0.77	0.77	231
provider	0.69	0.68	0.68	322
queries from insurance firm	1.00	0.33	0.50	21
queries from pharmacy	0.97	0.96	0.96	344
query on current appointment	0.52	0.20	0.29	130
refill	0.89	0.93	0.91	1933
rescheduling	0.67	0.59	0.62	304
running late to appointment	0.99	0.97	0.98	139
sharing of health records (fax, e-mail, etc.)	0.75	0.76	0.75	687
sharing of lab records (fax, e-mail, etc.)	0.76	0.60	0.67	277
symptoms	0.45	0.12	0.19	239
avg / total	0.80	0.81	0.80	10783

7.2. Multinomial Naïve Bayes:

The multinomial Naive Bayes classifier is suitable for classification with discrete features (e.g., word counts for text classification). The multinomial distribution normally requires integer feature counts. However, in practice, fractional counts such as tf-idf may also work.

I am creating pipeline for this using Countvectorizer and Tf-idf:

```
cat_text_NB = Pipeline([['vect', CountVectorizer(ngram_range=(1,2))],['tfidf', TfidfTransformer()],['clf',
MultinomialNB(alpha=1e-2)],])
```

The data is then divided into training and testing data in proportion of 80:20 using stratified sampling.

Classifier is trained and saved on harddisk:

```
with open('category_NB.pkl', 'wb') as f:
    pickle.dump(cat_text_NB, f)
```

Here we achieve accuracy of 77.3%

Now for predicting sub categories again we split data into training and test data in proportion of 80:20. Classifier of sub category prediction is made:

```
sub_cat_NB = Pipeline([['vect', CountVectorizer()],['tfidf', TfidfTransformer()],['clf', MultinomialNB()],])
```

Model is trained and saved as:

```
with open('sub_category_NB.pkl', 'wb') as f:
    pickle.dump(sub_cat_NB, f)
```

For predicting sub categories we predict categories first of our test data and use that also to predict our sub categories. The accuracy achieved here is 54.9% almost 55% which is quiet low.

8. Conclusion:

In the 2 models we prepared, our 1st model is quiet good as compared to our 2nd model so SGD classifier is very good in predicting categories and sub categories with respect to Multinomial Naïve Bayes. SGD classifier is giving around 82% accuracy in prediction of categories and 81% accuracy in prediction of sub categories.

9. Recommendations:

We have seen that our SGD classifier failed to predict sub categories like change of pharmacy and change of hospital. We need more data for these labels or we can use techniques like SMOTE (Synthetic Minority Over-Sampling Technique) for getting a more balanced dataset. We can study values of DATA and SUMMARY columns more carefully to remove words which are repeated almost in each row like call. Better the data better will be the prediction.